

HarvardX. Module 9: Data Science

Final Project: E-commerce conversion prediction

Adelaida Fernández Sanz

9/02/2020

Introduction and Aim of the Project

This project is part of the HarvardX Data Science Professional Certification: Capstone. The objective of the model is to be able to predict whether a user, defining user as a visitor to a certain e-commerce website, will make a purchase on the website or not. It is about predicting the conversion per user of the e-commerce. For this, the data obtained in the realization of the article Sakar, C.O., Polat, S.O., Katircioglu, M., Neural Comput & Applic (2018) will be taken as a basis. The database consists of vectorized variables that belong to 12,330 user sessions. Each session belongs to a different user obtained during a period of one year in order to avoid any tendency to a specific campaign, specific day, user profile or specific period.

Method & Analysis

We will apply two different algorithms: SVM (Linear) and Random Forest since they are very efficient in predictive tasks that require regression and classification techniques.

Firstly, the database has been inspected to ensure that there is no missing value and to identify and inspect the variables of the database. According to the source cited above, the database consists of:

- 10 numerical variables and 8 categorical variables.
- The dependent variable, the one that we want to predict in the model is the variable “Revenue”:the effectiveness in the purchase. It is a dichotomous variable: TRUE (If you buy), FALSE (do not buy).

The database is not balanced. The TRUE class is a minority (1908) compared to the FALSE class (10,422). So we will measure the performance of the model with the original database and with a balanced database to compare results and select the model with the best fit. We will use the downsampling technique in order to balance the data set. Machine learning classifiers like SVM or Random Forest do not deal very well with unbalanced training datasets as they are sensitive to the proportions of the different classes. As a consequence, these algorithms tend to favor the class with the highest proportion of observations (known as the majority class), which can lead to biased accuracy metrics.

For the SVM approach, to achieve the best possible fit, under both models (unbalanced & balanced), we will look after the best C parameters: The C Parameter (margin), represents the complexity constant. Specifies whether the model should be more generalized or more specific. The higher the value of the parameter, the greater the specificity, but this can lead to an overfitting. We will therefore test with 10 values in this range $C = [0.01 : 0.2]$. For the Random Forest approach we will look after the best mtry parameter, it defines the number of variables randomly sampled as candidates at each split.

Analysis Steps

0. Download packages & Data
1. Data set Exploration
2. Data set partition
3. SVM models: 3.1 Unbalanced Data 3.1.Balanced Data
4. Random Fores 4.1.Unbalanced Data 4.1.Balanced Data

Results

0. Download packages & Data needed

```
##0.Dataset and Packages downloading
##Packages Download
if(!require(tidyverse)) install.packages("tidyverse", repos = "http://cran.us.r-project.org")
if(!require(caret)) install.packages("caret", repos = "http://cran.us.r-project.org")
if(!require(data.table)) install.packages("data.table", repos = "http://cran.us.r-project.org")

library(tidyverse)
library(caret)
library(lattice)
library(ggplot2)
library(data.table)
library(dplyr)
library("readr")

# Online Shoppers Purchasing Intention Dataset, csv downloaded via:
#http://archive.ics.uci.edu/ml/datasets/Online+Shoppers+Purchasing+Intention+Dataset
#http://archive.ics.uci.edu/ml/machine-learning-databases/00468/

#Creating data-set:
data<-read.csv("online_shoppers_intention (1).csv")
```

1. Dataset Exploration

Here we can see: 1. The data structure:

```
## 'data.frame': 12330 obs. of 18 variables:
## $ Administrative : int 0 0 0 0 0 0 1 0 0 ...
## $ Administrative_Duration: num 0 0 0 0 0 0 0 0 0 ...
## $ Informational : int 0 0 0 0 0 0 0 0 0 ...
## $ Informational_Duration : num 0 0 0 0 0 0 0 0 0 ...
## $ ProductRelated : int 1 2 1 2 10 19 1 0 2 3 ...
## $ ProductRelated_Duration: num 0 64 0 2.67 627.5 ...
## $ BounceRates : num 0.2 0 0.2 0.05 0.02 ...
## $ ExitRates : num 0.2 0.1 0.2 0.14 0.05 ...
## $ PageValues : num 0 0 0 0 0 0 0 0 0 ...
## $ SpecialDay : num 0 0 0 0 0 0 0.4 0 0.8 0.4 ...
## $ Month : chr "Feb" "Feb" "Feb" "Feb" ...
## $ OperatingSystems : int 1 2 4 3 3 2 2 1 2 2 ...
```

```
## $ Browser          : int  1 2 1 2 3 2 4 2 2 4 ...
## $ Region           : int  1 1 9 2 1 1 3 1 2 1 ...
## $ TrafficType      : int  1 2 3 4 4 3 3 5 3 2 ...
## $ VisitorType      : chr  "Returning_Visitor" "Returning_Visitor" "Returning_Visitor" "Return
## $ Weekend          : logi  FALSE FALSE FALSE FALSE TRUE FALSE ...
## $ Revenue          : logi  FALSE FALSE FALSE FALSE FALSE FALSE ...
```

2. The most relevant statistics:

```
## Administrative      Administrative_Duration Informational
## Min.   : 0.000      Min.   : 0.00      Min.   : 0.0000
## 1st Qu.: 0.000      1st Qu.: 0.00      1st Qu.: 0.0000
## Median : 1.000      Median : 7.50      Median : 0.0000
## Mean   : 2.315      Mean   : 80.82      Mean   : 0.5036
## 3rd Qu.: 4.000      3rd Qu.: 93.26      3rd Qu.: 0.0000
## Max.   :27.000      Max.   :3398.75      Max.   :24.0000
## Informational_Duration ProductRelated      ProductRelated_Duration
## Min.   : 0.00      Min.   : 0.00      Min.   : 0.0
## 1st Qu.: 0.00      1st Qu.: 7.00      1st Qu.: 184.1
## Median : 0.00      Median : 18.00      Median : 598.9
## Mean   : 34.47      Mean   : 31.73      Mean   : 1194.8
## 3rd Qu.: 0.00      3rd Qu.: 38.00      3rd Qu.: 1464.2
## Max.   :2549.38      Max.   :705.00      Max.   :63973.5
## BounceRates          ExitRates          PageValues          SpecialDay
## Min.   :0.000000      Min.   :0.00000      Min.   : 0.000      Min.   :0.00000
## 1st Qu.:0.000000      1st Qu.:0.01429      1st Qu.: 0.000      1st Qu.:0.00000
## Median :0.003112      Median :0.02516      Median : 0.000      Median :0.00000
## Mean   :0.022191      Mean   :0.04307      Mean   : 5.889      Mean   :0.06143
## 3rd Qu.:0.016813      3rd Qu.:0.05000      3rd Qu.: 0.000      3rd Qu.:0.00000
## Max.   :0.200000      Max.   :0.20000      Max.   :361.764      Max.   :1.00000
## Month                OperatingSystems      Browser          Region
## Length:12330          Min.   :1.000      Min.   : 1.000      Min.   :1.000
## Class :character      1st Qu.:2.000      1st Qu.: 2.000      1st Qu.:1.000
## Mode  :character      Median :2.000      Median : 2.000      Median :3.000
##                      Mean   :2.124      Mean   : 2.357      Mean   :3.147
##                      3rd Qu.:3.000      3rd Qu.: 2.000      3rd Qu.:4.000
##                      Max.   :8.000      Max.   :13.000      Max.   :9.000
## TrafficType          VisitorType          Weekend          Revenue
## Min.   : 1.00      Length:12330      Mode :logical      Mode :logical
## 1st Qu.: 2.00      Class :character  FALSE:9462          FALSE:10422
## Median : 2.00      Mode  :character  TRUE :2868           TRUE :1908
## Mean   : 4.07
## 3rd Qu.: 4.00
## Max.   :20.00
```

3. The number of completed purchases (TRUE) and the number of not completed purchased (FALSE)

```
##
## FALSE  TRUE
## 10422  1908
```

As seen in the basic explorarion of the dataset, there is no missing values and modification of the Revenue feature class is needed:

```
#From logical to chr : Revenue
data <- data %>% mutate(Revenue = replace(Revenue, Revenue == "FALSE", "KO"))
data <- data %>% mutate(Revenue = replace(Revenue, Revenue == "TRUE", "GOOD"))
class(data$Revenue)
```

```
## [1] "character"
```

2. Dataset partition: Trainig & Test + Blanaced Dataset

First we partition the data set: 85% Training, 15% test. Secondly we prapared the balanced data set with the downsampling method.

```
##2.1Creating training-test dataset
n=nrow(data)
ind=1:n
itraining=sample(ind,floor(n*0.85))
itest=sample(setdiff(ind,itraining),floor(n*0.15))
training = data[itraining,]
testing = data[itest,]
dim(training)
```

```
## [1] 10480    18
```

```
dim(testing)
```

```
## [1] 1849    18
```

```
##2.2 Creating balanced dataset: Downsampling method
downSampled_training = downSample(x=training[, -ncol(training)],
                                   y=as.factor(training$Revenue))
downSampled_testing = downSample(x=testing[, -ncol(testing)],
                                  y=as.factor(testing$Revenue))
```

3. SVM Model


Let's start applying our SVM model. Firstly with unbalanced data, secondly with balanced data. The predictors features included in our model are the folowing:

```
predictors = names(training)[names(training) != "Revenue"]
predictors
```

```
## [1] "Administrative"      "Administrative_Duration"
## [3] "Informational"       "Informational_Duration"
## [5] "ProductRelated"     "ProductRelated_Duration"
## [7] "BounceRates"        "ExitRates"
## [9] "PageValues"         "SpecialDay"
## [11] "Month"              "OperatingSystems"
## [13] "Browser"            "Region"
## [15] "TrafficType"        "VisitorType"
## [17] "Weekend"
```

```
##3.1. Unbalanced Data: Training svm unbalanced data. Fiting the model-> C parameter
train_control<-trainControl(method="cv", number = 10, p = .9)
svm_fit_u_1 <- train(Revenue ~ ., method = "svmLinear", data = training,
                    trControl = train_control,
                    tuneGrid = expand.grid(C = seq(0.01, 0.2, length = 10)))
```

3.1 SVM Linear - Unbalanced Model In this plot we can see how well each C parameter performs:



```
# The best tuning parameter C that maximizes model accuracy
c<-svm_fit_u_1$bestTune
results_acc_by_c_u_1<-as_tibble(svm_fit_u_1$results[which.max(svm_fit_u_1$results[,2]),])
results_acc_by_c_u_1
```

```
## # A tibble: 1 x 5
##       C Accuracy Kappa AccuracySD KappaSD
##   <dbl>   <dbl> <dbl>       <dbl>   <dbl>
## 1 0.0311     0.887 0.458     0.00521 0.0278
```

```
# Applying best C
svm_def_u_1 <- train(Revenue ~ ., method = "svmLinear", data = training,
                    trControl = train_control,
                    cost= c)
```

```
## Overall Accuracy=0.88 , Sensitivity = 0.368, Specificity = 0.978,F = 0.859, PPV = 0.766
## , NPV = 0.888
```


```
##           Reference
## Prediction GOOD  KO
##      GOOD  111   34
##      KO    191 1513
```

As we can see the accuracy of the balanced SVM linear model is high, but the Sensitivity is very low. The specificity, of course, is very high taking into account that we are dealing with unbalanced data. As we can see the PPV is, of course, lower than the NPV.

##3.2. Balanced Data:SVM

```
##Training svm balanced data. Fiting the model-> C parameter
svm_fit_b_l <- train(Class ~ ., method = "svmLinear", data = downSampled_training,
                    trControl = train_control,
                    tuneGrid = expand.grid(C = seq(0.01, 0.2, length = 10)))
```

3.2 SVM Linear - Balanced Model In this plot we can see how well each C parameter performs:



```
# The best tuning parameter C that maximizes model accuracy
cbl<-svm_fit_b_l$bestTune
results_acc_by_c_b_l<-as_tibble(svm_fit_b_l$results[which.max(svm_fit_b_l$results[,2]),])
results_acc_by_c_b_l
```

```
## # A tibble: 1 x 5
##       C Accuracy Kappa AccuracySD KappaSD
##   <dbl>   <dbl> <dbl>      <dbl>   <dbl>
## 1  0.2    0.830 0.660    0.0188  0.0377
```

```
# Applying best C
svm_def_b_l <- train(Class ~ ., method = "svmLinear", data = downSampled_training,
  trControl = train_control, cost=cbl)
```

```
## Overall Accuracy=0.82 , Sensitivity = 0.752, Specificity = 0.887,F = 0.878, PPV = 0.870
## , NPV = 0.781
```

```
##           Reference
## Prediction GOOD  KO
##           GOOD  227 34
##           KO    75 268
```


Comparing to the unbalanced model, we can see how the overall accuracy has decrease, but the sensitivity has been balanced. In this case, the PPV is greater than the NPV.

4. RF Model

Now let's take a look into the RF results.

```
##3.2 Random Forest
##3.2.1. Training RF unbalanced data.
train_control<-trainControl(method="cv", number = 10, p = .9)
svm_fit_u_r <- train(Revenue ~ ., method = "rf", data = training, tuneLength=5,
  trControl = train_control)
```

4.1 Random Forest - Unbalanced Model In this plot we can see how well each mtry parameter



performs:

```
# The best tuning mtry that maximizes model accuracy
mtry_u<-svm_fit_u_r$bestTune
results_acc_mtry_u_r<-as_tibble(svm_fit_u_r$results[which.max(svm_fit_u_r$results[,2]),])
results_acc_mtry_u_r
```

```
## # A tibble: 1 x 5
##   mtry Accuracy Kappa AccuracySD KappaSD
##   <dbl>   <dbl> <dbl>       <dbl>   <dbl>
## 1     8     0.903 0.595     0.00855  0.0398
```

```
# Applying best mtry
svm_def_u_r <- train(Revenue ~ ., method = "rf", data = training, minNode=mtry_u$mtry,trControl
```

```
## Overall Accuracy=0.90 , Sensitivity = 0.586, Specificity = 0.966,F = 0.857, PPV = 0.770
## , NPV = 0.923
```


```
##           Reference
## Prediction GOOD   KO
##           GOOD 177   53
##           KO   125 1494
```

As we can see we have increase the Overall Accuracy comparing to the SVM unbalanced. Regarding Sensitivity , the unbalanced impact is not that significant compared to the SVM model. We are dealing again with unbalanced data so, as in the SVM model, the PPV is,of course lower than the NPV.

##3.2. Training RF balanced data.

```
##Training svm balanced data. Fiting the model-> C parameter
svm_fit_b_r <- train(Class ~ ., method = "rf", data = downSampled_training,
                    tuneLength=5,trControl = train_control)
```


4.2 Random Forest - Balanced Model In this plot we can see how well each mtry parameter performs:



```
# The best tuning mtry that maximizes model accuracy
cbr<-svm_fit_b_r$bestTune
results_acc_by_c_rb<-as_tibble(svm_fit_b_r$results[which.max(svm_fit_b_r$results[,2]),])
results_acc_by_c_rb
```

```
## # A tibble: 1 x 5
##   mtry Accuracy Kappa AccuracySD KappaSD
##   <dbl>   <dbl> <dbl>      <dbl>   <dbl>
## 1     2     0.856 0.713      0.0202  0.0405
```

```
# Applying best mtry
svm_def_b_r <- train(Class ~ ., method = "rf", data = downSampled_training,
                     minNode=cbr$mtry,
                     trControl = train_control,)
```

```
## Overall Accuracy=0.85 , Sensitivity = 0.864, Specificity = 0.831,F = 0.834, PPV = 0.837
## , NPV = 0.860
```

```
##           Reference
## Prediction GOOD  KO
##      GOOD  261  51
##      KO    41  251
```

The last model, presents higher level of Accuracy than the SVM balanced data and good balanced values of NPV and PPV.

Conclusion & Discusion

The following table presents all the relevant parameters of the four applied models

Model	Accuracy	Sensitivity	Specificity	TNrate	TPrate	F__measure
SVM-Unbalanced	0.8783126	0.3675497	0.9780220	0.8879108	0.7655172	0.8588194
SVM-Balanced	0.8195364	0.7516556	0.8874172	0.7813411	0.8697318	0.8784855
RF-Unbalanced	0.9037317	0.5860927	0.9657401	0.9227918	0.7695652	0.8565640
RF-Balanced	0.8476821	0.8642384	0.8311258	0.8595890	0.8365385	0.8338234

As we can see, the model with the highest accuracy is the RF with the original database (unbalanced) with a 90.5% accuracy. On the other hand, since it is an unbalanced database, it presents low sensitivity ratio and a low PPV ratio. The RF model with balanced bbdd has an acceptable accuracy (82%) and with very similar PPV and NPV ratios. Thus we can conclude that: RF in general has a better performance than SVM in both balanced and unbalanced databases.

Thanks for reviewing my work.

Adelaida Fernández